

Promptworks

Prompt engineering practice platform - website wireframes

B2B-led, free individual tier. Working name; see brief for alternates.

Design reference, August 2026

MARKETING

1. Global header + footer
2. Homepage
3. Platform / how it works
4. Tracks + challenge library
5. For Teams (the money page)
6. Pricing
7. Resources / prompt library
8. About + Contact / demo

APP (authenticated)

9. Onboarding
10. Learner dashboard
11. Exercise player <- core screen
12. Score + feedback
13. Workflows library
14. Team admin dashboard
15. Footer reference

Row 1 - utility bar / Row 2 - main row

Why prompt training?

For business

Docs

English | Espanol

Log in

logo

Promptworks

Platform

Tracks

For Teams

Pricing

Resources

Start free

Book a demo

App header (authenticated - different from marketing header)

Promptworks

Dashboard

Practice

Workflows

Team

12 day streak

Lvl 4

Footer - 4 columns, every page

Promptworks

Prompt engineering, turned from a training video into a measurable skill.

Product

Platform

Tracks

Workflows

Pricing

Learn

Prompt library

Guides

Blog

Certification

Company

About

Contact

Security

Book a demo

Notes

- No search in the marketing header - search lives in the app, scoped to the challenge library.
- Two CTAs always paired: 'Start free' (bordered, self-serve) and 'Book a demo' (filled, B2B). Never one without the other.
- Logo mark space reserved to the left of the wordmark, same as Carlos Tech Solutions.

Why prompt training?

For business

Docs

English | Espanol

Log in

logo

Promptworks

Platform

Tracks

For Teams

Pricing

Resources

Start free

Book a demo

Your team already has AI. Teach them how to ask.

Promptworks turns prompt engineering into reps - real work scenarios, scored feedback, and a dashboard that shows the skill actually moving.

Free for individuals. Team plans start where the ROI does.

Start practicing free

Book a demo

Scored in seconds by a fine-tuned grader

120+ role-based scenarios

SSO ready - your prompts never train our models

VALUE PROPOSITION

For the employee

Stop guessing. Get a score, a rewrite, and the reason it was better.

For the manager

See who is actually fluent - not who finished the video.

For the business

Fewer rewrites, fewer wasted tokens, fewer hours. Measured, not claimed.

The problem: the tools are on everyone's desk. The prompts are the bottleneck.

Short stat block - AI spend vs. measured behaviour change.

How it works

1. Pick a scenario

Real task from your role's track.

2. Write your prompt

Scenario left, prompt box right.

3. Get scored

Fine-tuned grader, 6 dimensions.

4. Ship it

Save the winner as a workflow.

TRY ONE NOW - live mini-exercise, no signup (highest-converting element on the page)

Scenario: deny a claim by email, empathetic, no liability, <150 words

[your prompt...]

Score it

Tracks by role

Sales

Service & Claims

HR

Marketing

Finance

Engineering

Homepage, lower half - continues from the role-track tiles

The rubric - six dimensions, identical everywhere in the product

Task clarity one line of plain english	Context one line of plain english	Constraints one line of plain english	Output format one line of plain english	Role & audience one line of plain english	Examples one line of plain english
--	---	---	---	---	--

ADVANCED TRACKS					
RAG & grounding		Agentic frameworks	Evals & data science	Fine-tune vs. prompt	Prompt injection & safety

AI-powered workflows

Practice is where prompts get good. Workflows are where they get used.
Save a validated prompt with variables, publish it to the team library, track usage.

workflow card screenshot

For Teams

Proficiency, adoption, hours returned - the dashboard screenshot that sells the product.

Median score
24/30

Weekly reps
312

Hours returned
~46

team heat map screenshot

Security & data handling

SSO, retention controls,
no training on customer data.

Case study slot

PLACEHOLDER - do not fabricate.
Leave empty until a real one exists.

Pricing teaser

Free / Pro / Teams / Enterprise.
No dollar figures yet.

Give your team ten minutes a week.

Start practicing free

Promptworks Prompt engineering, turned from a training video into a measurable skill.	Product Platform Tracks Workflows Pricing	Learn Prompt library Guides Blog Certification	Company About Contact Security Book a demo
---	--	---	---

Why prompt training?

For business

Docs

English | Espanol

Log in

logo

Promptworks

Platform

Tracks

For Teams

Pricing

Resources

Start free

Book a demo

One loop, repeated until it is a habit.

Scenario, prompt, score, rewrite, ship. Ten minutes at a time.

1

Pick a scenario

A real task from your role: a claims email, a discovery-call summary, a job description, a variance explanation. Constraints stated up front, the way a real brief states them.

screenshot

2

Write your prompt

Scenario and constraint checklist on the left, prompt box and model selector on the right. Token and cost estimate updates live - efficiency is part of the skill, not an afterthought.

screenshot

3

Get scored

The prompt runs against a live LLM. A grader model scores six rubric dimensions 0-5 and returns specific feedback: 'no output format given - the model guessed prose'.

screenshot

4

Compare and ship

Reference prompt side by side with yours, plus a diff. Save the better one as a reusable team workflow with variables. Practice output becomes a team asset.

screenshot

Why prompt training?

For business

Docs

English | Espanol

Log in

logo

Promptworks

Platform

Tracks

For Teams

Pricing

Resources

Start free

Book a demo

Tracks

Foundations first, then your role, then the advanced stuff. Preview any 3 scenarios free.

All

Foundations

By role

Advanced

Free preview

Search scenarios...

Sort: Difficulty

Foundations

6 scenarios . Beginner

Anatomy of a prompt, context, constraints, format, few-shot, iteration

Sales

14 scenarios . Mixed

Discovery summaries, objection handling, outreach that is not spam

Service & Claims

16 scenarios . Mixed

Denial explanations, escalation notes, tone under pressure

HR

12 scenarios . Mixed

Job descriptions, review summaries, policy answers with citations

Marketing

12 scenarios . Mixed

Briefs, variant generation, brand-voice constraint

Finance

10 scenarios . Advanced

Variance explanation, memo drafting, numbers you must not invent

Engineering

14 scenarios . Advanced

Spec to code, review prompts, test generation

Agentic frameworks

8 scenarios . Advanced

Multi-step plans, tool use, handoffs, failure recovery

Evals & data science

8 scenarios . Advanced

Writing a rubric, judging outputs, measuring drift

Promptworks

Dashboard

Practice

Workflows

Team

12 day streak

Lvl 4

Practice > Service & Claims > Denial explanation email

Scenario

A customer's claim was denied because the policy excludes flood damage. Write the prompt that produces the email to them.

Constraints

☐ Empathetic, plain english

☐ No admission of liability

☐ Under 150 words

☐ Cite the policy section

☐ End with one clear next step

Audience

Policyholder, upset, not an expert

Intermediate

Attempt 1 of unlimited

The reference prompt is hidden until you submit. That is the whole exercise.

Your prompt

Model: [claude-sonnet v]

[monospace prompt textarea]

0 / 2000 tokens

Est. 340 in / 180 out . ~\$0.004 per run

Run prompt

Model output

[streamed response appears here after Run]

Score panel opens below on completion (next page)

Promptworks

Dashboard

Practice

Workflows

Team

12 day streak

Lvl 4

22 / 30

Solid. Two dimensions are costing you most of the gap.

+140 XP . Level 4 . 12 day streak

Try again

Next scenario

Save as workflow

Rubric breakdown - same six dimensions on every screen in the product

Task clarity		5/5	Clear ask, no ambiguity about the deliverable.
Context supplied		4/5	Gave the denial reason. Missing: the customer already called twice.
Constraints		3/5	Word limit stated, but 'no admission of liability' was left implicit.
Output format		2/5	No format given - the model guessed prose when you wanted a subject line + body.
Role & audience		5/5	Named the reader and their state of mind. This is your strongest habit.
Examples (few-shot)		3/5	One example would have locked the tone. You gave none.

Your output

Reference prompt

Diff

Your prompt: 340 tokens . Reference: 210 tokens

[reference prompt shown side by side with yours - only after submission]

Cost efficiency is scored too. A prompt that gets the same answer in fewer tokens is a better prompt, and it is the cleanest ROI number the team dashboard has.

Promptworks

Dashboard

Practice

Workflows

Team

12 day streak

Lvl 4

Welcome back, Carlos

Service & Claims track . Level 4 . 12 day streak

Reps this week

7

goal 10

Median score

22/30

+3 this month

Weakest dimension

Output format

avg 2.4/5

Workflows shipped

3

used 41 times

Continue where you left off

Escalation note to a supervisor

Service & Claims . Intermediate . not started

Practice

Explain a coverage gap without jargon

Service & Claims . Intermediate . 1 attempt

Practice

Summarise a 40-minute call

Foundations . Beginner . completed 26/30

Practice

Your rubric profile

Task clarity		4.6
Context		3.9
Constraints		3.2
Output format		2.4
Role & audience		4.4
Examples		3.0

Output format is your gap. 3 scenarios queued for it.

 Promptworks

Dashboard

Practice

Workflows

Team

12 day streak

Lvl 4



Workflows

Prompts that scored well, saved with variables, shared with the team. This is where practice pays for itself.

All

Mine

Team

Service & Claims

Sales

Most used

+ New workflow

Workflow	Track	Inputs	Usage	Source score	Author
Claim denial email	Service & Claims	3 variables	used 41x	26/30	C. Salcedo
Discovery call summary	Sales	2 variables	used 33x	27/30	M. Reyes
Policy answer with citation	Service & Claims	4 variables	used 28x	25/30	A. Chen
Variance explanation memo	Finance	5 variables	used 12x	24/30	J. Patel
Job description draft	HR	3 variables	used 9x	23/30	L. Gomez

Workflow detail

Prompt template with {{variable}} slots filled by a small form. Non-experts run an expert's prompt without writing one.

Source attempt is linked - you can see how it was scored.

Why this section exists

Practice alone is a training product, and training products churn. Workflows make the skill produce something the org keeps using - which is the renewal argument and the ROI story at once. It is also where 'AI-powered workflows' stops being a buzzword and becomes a screen.

You bought the licences. Nobody can prove they worked.

Promptworks measures prompting skill the way you would measure any other one:
assign it, score it, watch the number move.

Book a demo

Assign

Push a track to a team or a person with a due date. Nudges by email.

Measure

Six-dimension heat map across the org. Find the gap, assign the fix in one click.

Benchmark

Your org against the anonymised platform median, by role. This is the data-science surface.

Ship

Winning prompts become shared workflows. Usage is tracked and reported.

Team admin dashboard preview

heat map screenshot - people x six rubric dimensions

Security

SSO / SAML, SCIM provisioning, retention controls, no training on your prompts or outputs.

Enterprise: end-to-end implementation

Scenarios authored from your own SOPs, a grader fine-tuned on your rubric, rollout and adoption reporting.

Proof

PLACEHOLDER for a real case study. Do not fabricate a logo wall or a customer quote in the mockup.

Team - Claims Operations

[Last 30 days v] [Export CSV]

Seats active

38 / 45

84% adoption

Median score

22/30

+3 vs. last month

Reps this week


312

8.2 per active user

Hours returned (est.)

~46

based on workflow usage



Weakest dimension

Output format - team avg 2.3 / 5

18 of 38 people below 3.

Assign 'Formatting outputs' track

Assignments

Foundations	Sep 12	31/45
Claims scenarios	Sep 26	12/45
Formatting outputs	not sent	-

Benchmark: your median 22 vs. platform median 19 for claims teams.

Free for one person. Priced for a team.

No dollar figures in the mockup - price after five customer conversations.

Free

\$0

- Foundations track

- 10 graded reps / month

- Community prompt library

- BYOK optional

Start free

Pro

\$ / month

- All tracks incl. advanced

- Unlimited graded reps

- Model comparison

- Certification

Upgrade

MOST TEAMS START HERE

Teams

\$ / seat / month

- Everything in Pro

- Admin dashboard + heat map

- Assignments & nudges

- Team workflow library

- SSO / SAML

Book a demo

Enterprise

Talk to us

- Custom scenarios from your SOPs

- Grader fine-tuned on your rubric

- End-to-end implementation

- Dedicated success manager

- SLA + security review

Book a demo

Open question: BYOK vs. you pay for inference on the free tier. BYOK removes your biggest cost risk but adds signup friction.

Resources / prompt library

Docs-style layout. Left sidebar of categories, right column of entries.

Foundations

Patterns

By role

Advanced

Safety

Prompt pattern entry - title, one-line use case, copy button

Prompt pattern entry - title, one-line use case, copy button

Prompt pattern entry - title, one-line use case, copy button

Prompt pattern entry - title, one-line use case, copy button

About

photo

Why this exists, in your own words. Short.

The technical background lives HERE, not on the homepage - LLMs, fine-tuned models, agentic frameworks, evals. Understated, recruiter-facing.

Contact / book a demo

Name / work email / company size / what you are trying to fix. Four fields, no more.

Direct email + calendar link.
Gainesville and Miami, FL.

Promptworks

Prompt engineering, turned from a training video into a measurable skill.

Product

Platform

Tracks

Workflows

Pricing

Learn

Prompt library

Guides

Blog

Certification

Company

About

Contact

Security

Book a demo

Onboarding - three screens, then straight into a rep

1. What do you do? Role picker: Sales, Service & Claims, HR, Marketing, Finance, Engineering, Other. Sets the default track. One question per screen.	2. What do you want to fix? Goal picker: write faster, get more consistent output, stop rewriting, train my team. Sets the first scenario.	3. Do one rep now Skip the tour. Drop them straight into a beginner scenario in their track. Signup completes AFTER the first score, not before.
---	--	--

Build order

STAGE 1	Frontend only - mock data, canned grading Homepage, Platform, For Teams, Pricing, exercise player, score screen, learner dashboard, team dashboard - all on mock JSON. Grading is a deterministic scorer that checks for constraint keywords. Ship to Vercel. Demoable in a weekend.
STAGE 2	Make it real FastAPI on Render + MongoDB Atlas. Live LLM calls behind a provider-agnostic adapter. Real rubric grading via LLM-as-judge returning strict JSON. Auth. The 'try one now' homepage widget wired to the real grader. Rate-limit POST /attempts hard.
STAGE 3	Sell it Team dashboard on real data, assignments, SSO/SAML, workflows library. Fine-tune a small grader model on your own accumulated stage-2 attempts - cheaper, faster, more consistent, and a real differentiator rather than a buzzword.

Decide before you open an editor: the name and domain, BYOK vs. paid inference on the free tier, and the first 15 scenarios. The scenarios are the moat and the real work. Everything else on this page is a weekend.